

FAILFORWARD AI

A Fan-Out and Adversarial-Gate Prompt Workflow for the Structured Autopsy of Failed Initiatives

FIELD	DETAIL
STUDENT NAME	K HARATHI
REGISTER NUMBER	99230040981
ASSESSMENT TYPE	Mini Assessment
WORKFLOW TYPE	Multi-stage prompt workflow (no code)
ARCHITECTURE	Fan-out → Convergence → Adversarial gate → Forward transform
STAGE COUNT	7 prompt stages + 1 conditional return prompt

1. TASK DESCRIPTION

This assessment required the design of a multi-step prompt workflow, built and executed entirely within an AI chat interface, without writing any code. The workflow had to break a single task into a sequence of clearly defined prompt stages, where the output produced at each stage is deliberately carried forward as input context for the next stage. The objective was to demonstrate structured prompt design and prompt chaining rather than to build a software application.

The workflow presented here satisfies that requirement and then departs from the usual shape in one important way. Instead of a single straight line of prompts, it uses a **fan-out, convergence and adversarial-gate architecture**: three stages run in deliberate isolation from one another on the same input, their outputs are then merged and their disagreements treated as evidence, and a later stage is permitted to send work backwards to be re-derived. The reasoning behind that choice is set out in Section 9.

2. CHOSEN PROBLEM AREA

Small initiatives fail quietly, and they fail constantly. A home-run subscription line closes after three months. A student club's mentorship programme fizzles out by week four. A first-time seller's product listing never converts. A neighbourhood tuition batch does not return after one term. In almost every one of these cases the people involved know that the initiative failed, but never learn with any precision why it failed.


What happens instead is that a single-sentence explanation is adopted within days — “there was no demand”, “the marketing was weak”, “people are simply not serious” — and then repeated until it becomes the accepted history of the event. That explanation is almost always the most visible symptom rather than the actual cause, and it is rarely tested against whatever evidence does exist. The real cost is not the initiative that failed; it is that the next attempt by the same person repeats the same structural mistake, because nothing transferable was ever extracted from the first one.

This workflow addresses that gap. It takes an initiative that has already ended or under-delivered, together with whatever evidence exists about it, and produces a defensible diagnosis of cause and a transferable prevention playbook that can be applied to the next initiative, not only to the one that failed.

The direction of this workflow is worth noting. It looks *backwards* at something that has already happened and converts it into forward-usable knowledge. It is a forensic workflow, not a planning one, and every design decision in it follows from that.

3. WORKFLOW ARCHITECTURE

The workflow consists of seven prompt stages arranged in four phases, with one conditional return path from the fourth phase back into the third.



Stages 2A, 2B and 2C are siblings, not successors. They receive identical input and are run without any knowledge of one another. Everything before them narrows the evidence; everything after them narrows the conclusion.

4. WORKFLOW STAGES AND EXPECTED OUTPUTS

- **Stage 1 — Evidence Ledger:** separates what is actually known from what is merely being assumed, before any cause is proposed. Every input fact is tagged OBSERVED (recorded or countable), REPORTED (someone said it) or ASSUMED (stated as fact with no basis given). This produces a numbered ledger of evidence items with their tags, plus an explicit list of the evidence that is missing. Every later stage must cite ledger item numbers.
- **Stage 2A — Numbers Lens:** reads only the countable evidence. This produces the shape of the failure over time — when decline began, how steep it was, whether it was a gradual slope or a sudden cliff, and what that timing alone rules in or out.

- **Stage 2B — People Lens:** reads only the human evidence. This produces the expectation gap — what participants or customers expected when they joined, what they actually received, and the precise point at which the two separated.
- **Stage 2C — Machinery Lens:** reads only the operational evidence. This produces an assessment of the machine behind the initiative — effort and cost per unit delivered, weekly time consumed, how many people it truly depended on, and the point at which operating load would have exceeded what was sustainable.
- **Stage 3 — Contradiction Merge:** merges the three lens outputs, marking where they agree, where they conflict, and where all three are silent. This produces a ranked cause hierarchy with exactly one root cause, separated from contributing factors and from coincidental factors that only look causal.
- **Stage 4 — Red-Team Challenge:** attacks the ranking against a fixed list of reasoning failures. This produces a verdict of SURVIVES, DOWNGRADE or REJECT for each cause with a named reason. If the root cause does not survive, the workflow returns to Stage 3 and re-ranks.
- **Stage 5 — Counterfactual Rebuild:** reconstructs what the initiative would have had to look like from day one for the surviving causes not to arise. This produces a rebuilt design, plus an honest verdict on whether the initiative was salvageable at all.
- **Stage 6 — Prevention Playbook:** converts each surviving cause into a pre-launch check and a measurable early-warning tripwire. This produces a one-page playbook usable by a person starting a different initiative who knows nothing about this one.

5. DETAILED PROMPT DESIGN

The prompts below were used, in the order shown, inside the AI chat interface. Stages 2A, 2B and 2C were each started in a fresh conversation containing only the Stage 1 ledger, so that none of them could see the others.

Stage 1 — Evidence Ledger

PROMPT 1

“Here is an initiative that has ended or under-delivered, together with everything I currently know about it: [insert description and all available evidence]. Do not diagnose anything yet. Build an evidence ledger. List each item of evidence as a numbered entry and tag it OBSERVED if it is a recorded number or countable fact, REPORTED if it is something a person said or believed, and ASSUMED if I have stated it as fact but have given you no basis for it. Then list separately the evidence that would normally exist for an initiative of this kind but that I have not provided. Do not propose a cause and do not comment on my own explanation.”

Stage 2A — Numbers Lens (isolated)

PROMPT 2A

“Using only the OBSERVED items in the ledger above, and ignoring every explanation I have offered, describe the shape of this failure over time. State when the decline began, how steep it was, whether it was a gradual slope or a sudden cliff, and whether it coincides with any dated event in the ledger. Then state what this timing alone makes unlikely. Cite ledger item numbers for every claim. Do not discuss people’s motives and do not discuss operational capacity.”

Stage 2B — People Lens (isolated)

PROMPT 2B

"Using only the ledger items that record what people said, did or expected, reconstruct the expectation gap. State what the participants or customers appear to have expected when they joined, what they actually received, and the specific point at which those two diverged. Where the ledger is silent about their experience, say that it is silent rather than filling the gap. Cite ledger item numbers. Do not discuss numbers, trends, timing or operational cost."

Stage 2C — Machinery Lens (isolated)

PROMPT 2C

"Using only the ledger items that describe how this initiative was actually run, assess the machinery behind it: the effort and cost per unit delivered, the time it consumed each week, how many people it genuinely depended on, and what happened to that load as volume grew. Identify the point at which operating load would have exceeded what the operator could sustain, whether or not demand had held. Cite ledger item numbers. Do not discuss customer motives and do not discuss timing trends."

Stage 3 — Contradiction Merge and Cause Ranking

PROMPT 3

"Below are three independent analyses of the same failure — a timing analysis, an expectation-gap analysis and an operational analysis — together with the evidence ledger all three were built from: [insert all four]. First, list every point on which two or more of them agree. Second, list every point on which they contradict each other, and state which reading the ledger better supports. Third, list what all three were unable to address. Then produce a single ranked cause hierarchy containing exactly one root cause, up to three contributing factors, and any coincidental factors that appear causal but are not. For each entry, cite the ledger items that support it and name which lens raised it."

Stage 4 — Red-Team Challenge

PROMPT 4

"Act as a hostile reviewer of the cause hierarchy above. Test every cause against the following list and name the specific test it fails: (a) hindsight bias, meaning it looks obvious only because we already know the outcome; (b) single-cause bias, meaning one cause is being credited with an effect larger than it could produce alone; (c) blame displacement, meaning fault is placed on customers or participants for behaving normally; (d) unfalsifiability, meaning no possible evidence could have disproved it; (e) weak anchoring, meaning it rests only on ASSUMED ledger items. Give each cause a verdict of SURVIVES, DOWNGRADE or REJECT with the reason. Do not propose replacement causes and do not soften the verdicts."

Stage 4R — Conditional Return (used only if the root cause does not survive)

PROMPT 4R

"The root cause has been [downgraded / rejected] for the following reason: [insert verdict]. Return to the cause hierarchy and re-rank it, treating the red-team verdict as new evidence rather than as an objection to be answered. You may promote a contributing factor to root cause, or you may conclude that the ledger does not support naming a root cause at all. Do not defend the previous ranking."

Stage 5 — Counterfactual Rebuild

PROMPT 5

"Using only the causes that survived the red-team gate, rebuild this initiative as it would have had to be designed from day one for those causes not to arise. State each change and the specific surviving cause it removes. Then answer directly: given those causes, was this initiative salvageable through redesign, or was it structurally unable to work as originally conceived? If it was not salvageable, say so plainly instead of producing a rescue plan."

Stage 6 — Prevention Playbook and Tripwires

PROMPT 6

"Convert the surviving causes and the rebuild into a prevention playbook for a person about to start a different initiative, who knows nothing about this one. For each surviving cause, give one pre-launch check written as a question they must be able to answer before starting, and one early-warning tripwire consisting of a specific measurable signal, a threshold, and the action to take when that threshold is crossed. Keep the playbook general enough to apply to a different initiative in a different field, and do not include any conclusion that did not survive Stage 4."

6. PROMPT CHAINING AND INFORMATION FLOW

In this workflow, deciding what each stage may not see is as much a design decision as deciding what it receives. The table below is the context ledger: it records, for every stage, what was carried forward, what was deliberately withheld, and why.

STAGE	RECEIVES	WITHHELD	REASON FOR WITHHOLDING
1 Ledger	Raw description and all evidence	Nothing	This stage is the anchor; it must see everything, including the operator's own explanation, in order to tag it as an assumption.
2A Numbers	Ledger only	The operator's narrative; Stages 2B and 2C	Prevents the model from reverse-fitting the numbers to a story it has already been handed.
2B People	Ledger only	Stage 2A output	Stops the expectation analysis from bending itself to explain a timing curve it has already seen.

STAGE	RECEIVES	WITHHELD	REASON FOR WITHHOLDING
2C Machinery	Ledger only	Stages 2A and 2B output	Keeps operational strain visible even when weak demand looks like the obvious answer.
3 Merge	Ledger and all three lens outputs	Nothing	Contradiction between the lenses is the entire purpose of this stage.
4 Red-team	Ledger and the cause ranking	The lenses' internal reasoning	Forces the ranking to defend itself on its stated evidence alone, not on how confidently it was derived.
5 Rebuild	Surviving causes only	Rejected and downgraded causes	Prevents a cause that failed the gate from re-entering the workflow through the redesign.
6 Playbook	Rebuild and surviving causes	The full case history	Forces the output to be a transferable rule rather than a story about this one initiative.

The single most important design decision in this workflow is what Stages 2B and 2C are *not* allowed to see. Each of the three lenses is an independent witness to the same event. The moment one of them reads another's testimony, it stops being a witness and becomes a corroborator, and the workflow loses the only thing that makes disagreement between them meaningful.

The chain is still a chain in every other respect. Stage 3 cannot run until all three lenses have reported. Stage 4 has nothing to attack until Stage 3 has produced a ranking. Stage 5 is defined entirely by what Stage 4 left standing, and Stage 6 cannot introduce a conclusion that Stage 4 rejected. The workflow is also auditable in both directions: any sentence in the final playbook can be traced back through the rebuild, the gate and the merge to a numbered ledger item, and any ledger item can be followed forward to see which lens used it and whether its conclusion survived.

7. DEMONSTRATION EXAMPLE

Initiative under autopsy. CrumbCrate, a weekly home-baking subscription box run by a single operator. It launched with sixty subscribers and had eighteen by the end of month three.

The operator's own explanation. "People lost interest, and a cheaper competitor opened nearby."

Evidence supplied. Weekly subscriber counts; delivery-day records; four cancellation messages; a note that each Saturday's boxes took about eleven hours to prepare; a price increase in week six; the operator's belief about the competitor.

Expected progression through the workflow:

- **Stage 1** tags the competitor claim as ASSUMED, since no cancellation message mentions it and no date is attached to it; tags the eleven-hour Saturday as REPORTED; and flags two missing items — there is no record of which week each cancellation fell in, and no log of delivery times.
- **Stage 2A** finds the losses were not gradual at all: roughly two-thirds occurred across weeks six and seven, immediately after the price change. A cliff, not a slope, which makes slow decay of interest unlikely.

- **Stage 2B**, having seen none of that, finds in the cancellation messages that subscribers expected variety and were receiving a repeating rotation by week four. The gap is between “a surprise each week” and “the same three items”.
- **Stage 2C**, having seen neither, finds that eleven unpaid hours for sixty boxes forced the menu to simplify as volume grew. The operating load was what produced the repetition.
- **Stage 3** notes that all three lenses converge on a week-six inflection but disagree about its cause, and ranks operator capacity forcing menu repetition as the root cause, the price rise as a contributing trigger, and the competitor as coincidental.
- **Stage 4** challenges the root cause for single-cause bias: capacity strain builds slowly and cannot by itself produce a two-week cliff. The cause is downgraded, and the return prompt re-ranks it as an eroded value proposition that had already been created by capacity strain, with the price rise as the trigger that finally made it visible to subscribers.
- **Stage 5** rebuilds the box around a capacity-honest menu of eight items batch-produced across two days, and states plainly that the original design could never have held sixty boxes inside one person's Saturday.
- **Stage 6** produces the pre-launch check “how many units can one person deliver at full quality in the hours actually available, and what happens at twice that number?”, with the tripwire “if the same item appears in consecutive deliveries more than once in four weeks, capacity has already been breached — cap new subscribers before cutting the menu.”

The value of the architecture is visible at Stage 4. Without the adversarial gate the workflow would have shipped a confident and wrong root cause, and the playbook would have taught the next operator to watch the wrong signal. The figures used above are illustrative and are included to show the shape of the workflow, not as data from a real business.

8. WHAT THE WORKFLOW IS ALLOWED TO CONCLUDE

Several prompts in this workflow are written to permit an answer the user would not want. Stage 4 may reject the root cause outright. Stage 4R may conclude that the evidence supports no root cause at all. Stage 5 may state that the initiative was never salvageable rather than produce a rescue plan. Stage 1 may tag the operator's own explanation as an assumption in the first output the user ever sees.

This is deliberate. A workflow whose every stage can only produce an encouraging answer is not a diagnostic instrument, because its output carries no information: it would have said the same thing regardless of the evidence. Making failure a permitted output at three separate stages is what allows a positive result from this workflow to mean anything.

9. WHY THIS WAS NOT BUILT AS A SINGLE PROMPT, OR AS A STRAIGHT CHAIN

Two comparisons are relevant here, not one.

Against a single prompt. A single broad prompt asking an AI why an initiative failed produces a fluent narrative that accepts the operator's own framing, because that framing is the loudest thing in the prompt. It will not separate evidence from assumption, it will not surface a cause that contradicts the story it was handed, and it offers no way to tell which part of its answer is load-bearing.

Against a straight chain. A linear chain solves the first problem but introduces a subtler one: contamination. In a straight chain, whatever the second stage concludes becomes the frame the third stage reasons inside. If the timing analysis runs first and reports a cliff at week six, a following stage asked about customer expectations will tend to explain that cliff rather than independently examine what customers expected. Its analysis will look like agreement, but it is inheritance. Three isolated lenses produce three genuinely independent readings, and where they disagree, that disagreement is diagnostic information which a chain would have destroyed before it was ever visible.

The adversarial gate exists for the same reason in the opposite direction. Every stage up to Stage 3 is building a case, and each one inherits the previous stage's confidence along with its content. Stage 4 is the only stage with no stake in the conclusion, and it is given a fixed list of named failure tests rather than being asked for a general opinion, so that its verdict is itself checkable. The conditional return from Stage 4 to Stage 3 is what makes this a workflow rather than a pipeline: an output can be sent back and re-derived, which a straight chain has no mechanism to do.

PROPERTY	SINGLE PROMPT	LINEAR CHAIN	THIS WORKFLOW
Evidence discipline	None; assumption and fact are blended	Possible if requested	Enforced at Stage 1 and cited at every later stage
Independence of views	None	Low; each stage inherits the last one's frame	High; three lenses run blind to each other
Self-criticism	None	Only if a critique stage is added	Built in as a gate the ranking must pass
Can revise itself	No	No; output only moves forward	Yes; one conditional return path
Traceability	None	Stage to stage	Every final claim traces to a numbered ledger item
Transferable output	Rarely	Sometimes	Required; Stage 6 is withheld the case history

10. LIMITATIONS AND FAILURE MODES OF THE WORKFLOW ITSELF

- **Garbage anchoring.** Stage 1 can only tag the evidence it is given. An assumption written confidently enough to resemble a recorded fact may still be tagged OBSERVED. The ledger reduces this risk; it does not remove it.
- **Loop abuse.** The return path from Stage 4 can be run repeatedly until a comfortable answer appears. In practice the return should be taken at most once, and a second failure at the gate should be recorded as “no root cause supportable on this evidence” rather than re-run.
- **Isolation is manual.** Because this runs in a chat interface, isolating Stages 2A to 2C requires starting each in a fresh conversation with only the ledger pasted in. Running all three in one thread silently collapses the fan-out back into a chain and removes the workflow's main advantage, while still appearing to work.
- **Thin-evidence cases.** With fewer than roughly ten ledger items the three lenses will largely repeat one another, and the cost of the fan-out is not repaid. Below that threshold a linear chain is the more honest choice.

11. EXPECTED LEARNING OUTCOME

This exercise built practical experience in treating a prompt sequence as an architecture rather than as a list. The central skill it developed was deciding, at each step, not only what information to carry forward but what to deliberately hold back, and being able to justify each withholding — the fan-out only works because three stages are kept ignorant of one another on purpose.

It also required writing prompts that constrain the shape of an answer rather than requesting a topic: forcing every claim to cite a numbered evidence item, forcing a verdict from a fixed set of three words, forbidding a critique stage from proposing replacements so that it cannot quietly become an author. A further lesson was that a workflow becomes trustworthy only when it is allowed to return an unwelcome result, which is why three separate stages here can conclude that there is nothing to salvage.

Finally, it showed that a back-edge — a single permitted return from a later stage to an earlier one — changes the character of a prompt sequence entirely, turning a set of instructions that can only accumulate confidence into one that can withdraw a conclusion it has already drawn.

12. CONCLUSION

This workflow demonstrates that a structured sequence of prompts, executed entirely inside a conversational AI interface and without a line of code, can take an initiative that has already failed and carry it through evidence anchoring, three independent diagnostic readings, contradiction merging, adversarial challenge, counterfactual redesign and finally a transferable prevention playbook.

Its distinguishing feature is not its length but its shape. By running three stages in isolation, treating their disagreement as evidence, subjecting the merged result to a stage designed to attack it, and permitting a conclusion to be sent back and re-derived, it produces a diagnosis that is harder to reach and easier to defend than anything a single prompt or a straight chain would return — and it does so purely through prompt design.